S0: BER of BPSK in AWGN

Chang Young Ham

September 2026 (accompanies `src/s0_bpsk_awgn_ber.py`)

Purpose. S0 is the calibration stage of the repository. S1 and S2 reuse its signal model ($r = s + n$, threshold at zero), so the simulator is first required to reproduce a result whose exact answer is known. This note derives that answer and the statistical test the script uses to decide whether simulation and theory agree. Mathematics used: Gaussian random variables, conditional probability, the binomial estimator.

1 Signal model

One bit $b \in \{0, 1\}$ is transmitted per symbol. Binary phase-shift keying maps it to

$$s = \begin{cases} +\sqrt{E_b} & b = 1 \\ -\sqrt{E_b} & b = 0 \end{cases}$$

so that every symbol carries energy $s^2 = E_b$ (energy per bit). The channel adds white Gaussian noise. After the receiver's matched filter the decision variable is a single real number

$$r = s + n, \qquad n \sim \mathcal{N}\big(0,\ \sigma^2\big)\,.$$

Where $\sigma^2 = N_0/2$ comes from. N_0 is the *one-sided* noise power spectral density, so the two-sided density is $N_0/2$. Passing white noise through a filter with impulse response $h(t)$ gives an output variance

$$\sigma^2 = \int_{-\infty}^{\infty} \frac{N_0}{2}\,|H(f)|^2\,df = \frac{N_0}{2}\int_{-\infty}^{\infty} |h(t)|^2\,dt = \frac{N_0}{2}$$

for a unit-energy matched filter (the middle step is Parseval's relation). Using $\sigma^2 = N_0$ instead shifts the whole simulated curve by exactly 3 dB.

2 The decision rule and the error event

Both symbols are equally likely, and the noise is symmetric, so the maximum-likelihood decision is a threshold at zero:

$$\hat{b} = \begin{cases} 1 & r > 0 \\ 0 & r \le 0. \end{cases}$$

Condition on $b = 1$ (so $s = +\sqrt{E_b}$). An error occurs when $r \le 0$, i.e. when $n \le -\sqrt{E_b}$:

$$P(\text{error} \mid b = 1) = P\Big(n \le -\sqrt{E_b}\Big) = P\left(\frac{n}{\sigma} \le -\frac{\sqrt{E_b}}{\sigma}\right) = \mathrm{Q}\left(\frac{\sqrt{E_b}}{\sigma}\right),$$

where $\mathrm{Q}(x) = P(Z > x)$ for $Z \sim \mathcal{N}(0,1)$ is the Gaussian tail function and we used the symmetry $P(Z \le -x) = P(Z \ge x)$. By the same argument $P(\text{error} \mid b = 0)$ is identical, so the average bit error rate is

$$\mathrm{BER} = \mathrm{Q}\left(\frac{\sqrt{E_b}}{\sigma}\right) = \mathrm{Q}\left(\sqrt{\frac{E_b}{N_0/2}}\right) = \boxed{\mathrm{Q}\left(\sqrt{\tfrac{2E_b}{N_0}}\right)}. \tag{1}$$

Numerical check. The last point of the S0 sweep is $E_b/N_0 = 9\,\mathrm{dB} = 7.94$. With $E_b = 1$, $\sigma = \sqrt{1/15.89} = 0.2509$ and the argument of Q is $1/0.2509 = 3.986$, so $\mathrm{BER} = \mathrm{Q}(3.986) = 3.36 \times 10^{-5}$, the value in `results/s0_ber_results.csv`. At $E_b/N_0 = 10\,\mathrm{dB}$, the thermal-noise condition used in S1, the same calculation gives $\mathrm{Q}(4.47) = 3.9 \times 10^{-6}$.

3 Q and erfc

The code uses `scipy.special.erfc` because SciPy does not provide Q directly. The two are related by

$$\mathrm{Q}(x) = \tfrac{1}{2}\,\mathrm{erfc}\left(\tfrac{x}{\sqrt{2}}\right) \quad \Longrightarrow \quad \mathrm{BER} = \tfrac{1}{2}\,\mathrm{erfc}\left(\sqrt{\tfrac{E_b}{N_0}}\right),$$

which is `ber_bpsk_theory()` in the script.

4 Where the model enters the code

The script fixes $E_b = 1$, so $N_0 = 1/(E_b/N_0)$ and $\sigma = \sqrt{N_0/2}$.

Mathematics	Line in `s0_bpsk_awgn_ber.py`
$\sigma = \sqrt{1/(2\,E_b/N_0)}$	`sigma = np.sqrt(1.0 / (2.0 * ebn0_lin))`
$r = s + n$, $s = \pm 1$	`received = symbols + noise`
$\hat{b} = [r > 0]$	`decisions = (received > 0).astype(int)`
Eq. (1)	`0.5 * erfc(np.sqrt(ebn0_linear))`
z-score, Eq. (2)	`(n_err - expected) / np.sqrt(expected)`

`scale` in `rng.normal` is the standard deviation σ, not the variance σ^2.

5 Accuracy of the Monte Carlo estimate and the pass criterion

The simulator observes K errors in n bits and reports $\widehat{\mathrm{BER}} = K/n$. Each bit is an independent Bernoulli trial with success probability $p = \mathrm{BER}$, so $K \sim \mathrm{Binomial}(n, p)$ and

$$\mathbb{E}[K] = np, \qquad \mathrm{Var}[K] = np(1-p) \approx np.$$

The *relative* standard deviation of the estimate is therefore $1/\sqrt{np} = 1/\sqrt{\mathbb{E}[K]}$.

Stopping rule. The script simulates in chunks of 10^6 bits and stops a point once at least `MIN_ERRORS = 200` errors have been seen. This bounds the relative spread by $1/\sqrt{200} \approx 7\%$ at the *hardest* point: at 9 dB the required $n \approx 200/p$ is six million bits. The easy points collect far more errors in their first chunk (79,271 at 0 dB, a spread of 0.36%), so the accuracy differs by a factor of twenty across the sweep.

Pass criterion. Because of that difference a single tolerance on the ratio $\widehat{\mathrm{BER}}/\mathrm{BER}$ would be either too loose at 0 dB or too tight at 9 dB. The script instead standardises each point,

$$z = \frac{K - np}{\sqrt{np}}, \tag{2}$$

which is approximately $\mathcal{N}(0,1)$ at every point if the simulator is correct, and requires $|z| < 4$. For a correct simulator the probability of failing at a given point is $P(|Z| > 4) = 6 \times 10^{-5}$, below 10^{-3} for the whole ten-point sweep. The observed values lie between -2.11 and $+2.22$, which is the expected scatter of ten standard normal draws. A modelling error is detected easily: $\sigma^2 = N_0$ instead of $N_0/2$ gives $z > 100$ at every point.

6 Link to S1

For large x, $\mathrm{Q}(x) \approx e^{-x^2/2}/(x\sqrt{2\pi})$, so with $x^2 = 2E_b/N_0$ the BER falls like e^{-E_b/N_0}: exponentially in the linear E_b/N_0. This exponential decay is the benchmark that partial-band jamming removes in S1, where the decay becomes only inverse-linear.